

Model loss over epochs (turbo gradient; dashed = validation) — best val_loss=0.0157 @ epoch 38

